

Web Bundles

Felipe Erias, Dan Ehrenberg
Igalia, in partnership with eyeo

WPACK - IETF 111

July 2021

Index

1. Proposed spec changes
2. Bundling and resource preloading
3. Bundling and the JS ecosystem

Proposed spec changes

Overview

- Goal: base Web Bundle spec as simple as possible
- Focus on consensus functionality
 - use cases like resource preloading
- Allow for extensions
 - use cases that require a more complex bundle format can be addressed in other documents

Remove the primary URL

<wpack-wg/bundled-responses/pull/5>

- Browsing inside a bundle requires a primary URL specified in the bundle
- Primary URL could be provided in a separate section
- However, it is outside the scope of the core bundle specification to define this section

Remove manifest section

wpack-wg/bundled-responses/pull/6

- As previous, outside the scope of the core spec
- Perhaps redundant?
 - primary URL response could already include it

```
Content-Security-Policy: manifest-src <source>;
```

Remove content negotiation

[wpack-wg/bundled-responses/pull/7](#)

- More efficient to do it on the server to select which bundle to send, rather than including several variants in the bundle and doing it on the client
- Related discussions:
 - [bundled-responses/issues/2](#)
 - [WICG/webpackage/pull/618](#)

Add bundle-format.md


[wpack-wg/bundled-responses/pull/8](https://github.com/wpack-wg/bundled-responses/pull/8)

- A summary of the specification, providing a formal description of the format of the bundle
 - no new information
- This PR assumes that all the others are accepted, so we can see their collective result

Updated bundle format

Name	Description
<code>magic</code>	<code>F0 9F 8C 90 F0 9F 93 A6</code> (🌐📦)
<code>version</code>	spec version, no change except as last resort
<code>section-lengths</code>	array of <code>section-length</code>
<code>sections</code>	array of sections
<code>length</code>	length of the bundle in bytes

Updated bundle format

<code>section-length</code>	<i>(updated)</i>
<code>section-name</code>	name of the section
<code>length</code>	length of the section in bytes
Sections 	
<code>index</code>	map from a URL to a <code>location-in-responses</code>
<code>critical</code>	<i>(optional)</i> sections that the client needs to understand
<code>responses</code>	array: HTTP responses for the bundle's representations

Updated bundle format

<code>location-in-responses</code>	<i>(updated)</i>
<code>offset</code>	Offset within the <code>responses</code> section
<code>length</code>	Size of the response in bytes
<code>response</code>	
<code>headers</code>	map of field names to field values
<code>payload</code>	content of the HTTP response as a byte string

Second document

[wpack-wg/bundled-responses/issues/9](#)

- Proposal: WPACK group creates a second document:
 - an extension focused on *"browsing into a bundle"*
- Additional sections may be specified in "critical"
- Goals:
 - simple, consensus base spec
 - extension spec: less general use cases that require a more complex bundle format

Discussion

- Pull requests
 - Remove the primary URL section
 - Remove the manifest section
 - Remove content negotiation
 - Add bundle-format.md
- Issue/proposal
 - Extension document for browsing into a bundle

Bundling and resource preloading

Proposal: [WICG/resource-bundles](#)

References

- Igalia: “**Bundle Preloading**”
 - [WICG/resource-bundles](#)
- Chromium has an experimental implementation of very similar ideas, this document enumerates the small differences between the two:
 - “*Web Bundles and Bundle Preloading*” [link](#)

Background

- Web sites are composed of multiple resources
 - hundreds or thousands in real-world web apps
 - fetching them one by one has poor performance
- *Bundlers*
 - combine and transform resources for deployment
 - **webpack**, **browserify**, **rollup**, **parcel**, **esbuild**...

Bundlers

- Dependency graph
- *Bundling*
 - combine resources for efficient deployment
- *Tree shaking*
 - identify and remove unused code
- *Virtualization*
 - include other resources inline in JS code
 - fewer network exchanges, but slower to load
- *Code splitting*
 - divide code/resources into various chunks
 - can be loaded on demand or in parallel

Motivations

- Efficiency
 - Retrieving bundled resources is more costly
 - Does not fit well with browsers' caching strategies
- Interoperability
 - Bundling strategies are not standardized
 - Bundling strategies are not interoperable

Goals

- Efficiently distribute web content
 - Web platform: standard, interoperable
- Keep benefits of today's bundler ecosystem
 - improve network performance, load content faster
 - facilitate revving, code splitting, tree shaking, etc.
 - reduce bundler build time and logic
- Keep benefits of accessing individual resources
 - flexibility when loading and processing responses
 - each response can be cached individually

Outline

- **Resource preloading with bundled responses**
 1. *document* provides a list of required resources
 2. *client* sends HTTP request for those not in its cache
 3. *server* replies with a bundled response
- **API components**
 - static API: declare resources to retrieve in a bundle
 - imperative API in JS: preload bundled resources
 - request headers, corresponding response behavior

Preserving

- Resource identity
 - bundles contain same resources as individual URLs
- Origin model
 - resources within the same origin as the bundle
- Path restriction
 - resources have the same `path-1` as the bundle
- URL consistency
 - each URL corresponds to a resource
 - same response for individual and bundled requests
 - can verify that servers are well-behaved
 - graceful degradation, content blocking

Usability goals

- For developers
 - no need to change tools (except bundlers)
 - code splitting, transformations, etc. remain possible
- For final and intermediary servers
 - no large time/space penalties if not supported
 - stateless serving logic
 - deploy static content just by copying files
- For browsers
 - fit current fetching and caching architectures
 - support graceful degradation

Privacy goals

- Personalization
 - do not enable disguising personalized content
 - do not weaken the significance of URLs.
- Content blocking
 - must be compatible with content blocking
 - "trusted" intermediary can not "repackage" sites
 - do not enable cheap rotation of URLs in the bundle
 - do not download blocked content

HTML `<script>` tag

```
<!-- https://www.example.com/index.html -->
<script type="bundlepreload">
  {
    "source": "./assets/resources.wbn",
    "resources": [
      "render.js",
      "profile.png"
    ]
  }
</script>
```


JavaScript API

```
// https://www.example.com/index.html
window.bundlePreload({
  source: "./assets/resources.wbn",
  resources: ["render.js", "profile.png"]
});

let image = document.createElement("img");
image.src = "assets/profile.png";
...
```

HTTP request

```
GET /assets/resources.wbn HTTP/1.1
...
Host: www.example.com
Bundle-Preload: "render.js", "profile.png"
...
```

HTTP response

- The response must be a *bundled response* containing HTTP responses for each of the requested URLs:

```
https://www.example.com/assets/render.js  
https://www.example.com/assets/profile.png
```

- These resources may be cached and references to them later on may be loaded from the cache

Code splitting

- `example.com/page1.html` lists resources `A`, `B`, `C`
 - client requests all these resources and caches them
- `example.com/page2.html` lists resources `A`, `D`, `E`
 - client only needs to request `D` and `E`
- As the client may request a subset of the listed resources, it is able to retrieve those it already has from its cache

Responsiveness

- Web devs may prioritize some resources to shorten the time until the page is ready for user interaction
 - e.g. declaratively, when the page is loaded
- Less urgent resources may be fetched after the page has completed its first load
 - e.g. imperatively, on `window.onload`

Avoid cachebusting

- Prevent the need to download a whole bundle when only some of its resources have changed
- Devs can update resources with fine granularity:
 1. append a version to the resource's name ("revving")
 2. update the preload resource list in the document

Summary

- API for **resource preloading with bundled responses**
 - declarative (HTML) and imperative (JavaScript)
 - specific request headers and response behaviour
- Preserves user privacy and freedom
- Backwards compatibility for servers, intermediaries
- Remove overhead from resource "virtualization", etc.
- No significant changes to dev workflow

Bundling and the JS ecosystem

JS Modules today

- JS APIs require modules to be in separate source files:
 - static `import` and dynamic `import()`
 - Web workers
 - worklets (paint, audio, animation, layout)
- This causes several problems:
 - APIs are more inconvenient to use
 - ergonomic solutions are insecure or non-standard
 - importing a source file from another origin affects relative paths, CORS, etc.

Goal of module fragments

- Allow bundling of multiple JS modules in a single file
 - can be used as output format for bundlers
 - also convenient to write for developers
- Complementary to resource bundle preloading
 - number of JS files tends to blow up in web apps
- Complementary to *code splitting*, *tree shaking*, etc.
 - optimize number and composition of JS source files
 - more efficient for chunks to be more granular, to provide exactly the ones needed

Standardization track

- This proposal at TC39:
 - **module fragments** (Stage 1)

Core idea

Declare a module inline, inside another source file:

```
module myModule {  
  export function myFunction(...) {  
    ...  
  }  
};
```

- Related proposal: **module blocks** (stage 2)

Importing

These modules can then be imported dynamically:

```
let myExports = await import(myModule);  
myExports.myFunction();
```

Or statically (from within other modules):

```
module combined {  
  import { myFunction } from myModule;  
  ...  
}
```

Exporting

Can be exported, so they are accessible from outside:

```
// Only accessible within this file
module priv {
  export function doInPrivate() { ... }
}

// Accessible from outside this file
export module pub {
  import { doInPrivate } from priv;
  export function doInPublic() { ... }
}
```

With resource preloading

- Resource bundle preloading
 - works at the network level
 - supports resources of any MIME type
 - supports additional metadata as HTTP headers
 - each resource can be cached individually
- JS module fragments
 - limited to JS source files
 - work at the JS module loading level
 - only affect how JavaScript is loaded
 - equivalent to what bundlers do today

Wrap up

- High-performance, standard bundling for the Web
 - reduce and reorganize the number of source files
 - distribute these and other resources efficiently
- Next steps:
 - module fragments is at Stage 1 in TC39
 - prototype module fragments together with bundle preloading

Contact

Felipe Erias felipeerias@igalia.com

Daniel Ehrenberg dehrenberg@igalia.com



Thank you!

